

White Paper: Bayesian Methods in Nuclear Physics

Edited by R.J. Furnstahl, D. Higdon, N. Schunck, and A.W. Steiner

June 27, 2016

Abstract

The INT program on “Bayesian Methods for Nuclear Physics”

Contributing Authors:

Dick Furnstahl

Department of Physics
Ohio State University

Dave Higdon

Biocomplexity Institute
Virginia Tech

Nicholas Schunck

Nuclear and Chemical Science Division
Lawrence Livermore National Laboratory

Andrew Steiner

Department of Physics and Astronomy
University of Tennessee / ORNL

Kyle Wendt

Institut für Kernphysik
Technische Universität Darmstadt

Contents

1	Overview	4
1.1	Goals	4
1.2	Building this document	5
1.3	Questions	6
2	Background	8
2.1	Basics of Bayesian statistics	8
2.2	Conceptual or philosophical issues	9
3	Themes	9
3.1	Implication of Bayesian philosophy/practices/concepts	9
3.2	Discrepant functions / random effects	9
3.3	Gaussian Process (GP) as a tool	9
3.4	χ^2 and dofs	9
3.5	Pitfalls	9
3.6	Model selection / comparison (or metrics)	9
3.7	Selection of priors	10
3.8	Use of MCMC	10
3.9	Specific issues	10
3.10	Understanding of model function	10
A	Guide to the literature (written and software)	11
A.1	General Bayesian statistics	11
A.2	Gaussian Process (GP) models	11
A.3	Hypothesis Testing and Model Comparison	11
B	Summary of notation and terminology	11
B.1	Lexicon	11
C	Using git and github	12
	References	13

1 Overview

- The INT program on “Bayesian Methods for Nuclear Physics” brought together statisticians and nuclear practitioners, principally theorists, to explore how Bayesian inference can enable progress on the frontiers of nuclear physics and open up new directions for the field.
- The program also served as ISNET-4, the fourth meeting in a series helping researchers bridge the gap between experiment and theory (ISNET stands for Information and Statistics in Nuclear Experiment and Theory).

With the maturation of calculational methods such as lattice QCD for hadronic physics, ab initio and density functional theory approaches for nuclear structure and reactions (with applications to astrophysics and fundamental symmetries), and viscous hydrodynamic modeling of relativistic heavy-ion collisions, nuclear theory is entering an era of precision calculations. This is leading to increased demand for sophisticated uncertainty quantification, to effectively interface with, inform, and analyze experiments. The methods used to quantify errors are often based on frequentist statistical analysis, but Bayesian methods are becoming increasingly popular.

Bayesian statistics is a well-developed field, although it has not been part of the traditional education of nuclear theorists. In schematic form, Bayesian statistics treats the parameters or the model/theory as genuine random variables. It then uses Bayes theorem of probabilities to provide a recipe to compute their probability distribution (the “posterior”) in terms of prior information (e.g., about the data) and a likelihood function. For applications to fitting (“parameter estimation”), the posterior lets us infer, given the data we have measured, the most probable values of the parameters and predict values of observables with confidence intervals. Other applications involve deciding between alternative explanations or parameterizations (“model selection”). In practice, there are pitfalls in the implementation of this formalism and it is often a computationally hard problem.

Interest in Bayesian statistics has increased significantly in the past 10 years. The wide availability of large-scale computing resources has made the computation of the integrals needed for Bayesian inference easier. Modern experimental and observational facilities generate large amounts of data, often best analyzed and characterized through Bayesian methods. Bayesian methods are often preferred for under-constrained fits and inverse convolutions. In nuclear science, Bayesian methods have found their way into such areas as nuclear data, lattice QCD, dense matter, effective field theory, nuclear reactions, and parton distribution functions. These sub-fields have generally turned to Bayesian inference methods independently and in some cases without access to expert advice and guidance from professional statisticians.

1.1 Goals

Among the goals of the program and this white paper are to

- facilitate cross communication, fertilization, and collaboration on Bayesian applications among the nuclear sub-fields;

- provide the opportunity for nuclear physicists who are unfamiliar with Bayesian methods to start applying them to new problems;
- learn from the experts about innovative and advanced uses of Bayesian statistics, and best practices in applying them;
- learn about advanced computational tools and methods;
- critically examine the application of Bayesian methods to particular physics problems in the various subfields;
- build a knowledge base.

1.2 Building this document

[This subsection is only part of the working document, not the final version.]

- Editors are the organizers of the program (Furnstahl, Higdon, Schunck, Steiner). In this capacity they are to provide a coherent narrative and make the final decisions on what content is included.
- Contributors are self-selected participants in the program. They might be designated the Bayes-INT Working Group.

Guidelines for building the document:

- Do not make firm decisions on the order/organization of the content until later, to make it flexible.
- Possible LaTeX formats:
 - We will build the document with a convenient structure (e.g., with a table of contents) and later decide on a publication plan.
 - The Journal of Physics G LaTeX format is a possibility because this journal is a likely target for a published version.
 - To get started, we will borrow a format from an existing white paper.
- Organization of files
 - Use Git, hosted on Github. Each contributor can make their own branch, and resolution of conflicts from merging will be made by the Editors.
 - Standardized list of macros in macros.tex. Use a macro whenever possible to facilitate developing a standardized notation, which can be easily modified.
 - Use BibTeX for references (newer version?). Include the titles in the style to more easily identify sources.
 - For early passes through new materials, use bullets (or enumerated lists) to enter content. This makes it modular and easy to adjust the ordering and structure.

- Figures
 - Plan to use original figures for schematic illustration as much as possible
 - But point specifically to figures in the literature
- Associated computer codes. We can include as part of the supplementary material Mathematica and Ipython notebooks.

1.3 Questions

In the following subsections are (partial) lists of questions that were considered during the INT program. [*Synthesize these into a more compact list of questions that we plan to address in this white paper.*]

1.3.1 General questions

- What do Bayesian techniques offer that frequentist statistics do not?
 - Also, what kinds of problems ill-suited for Bayesian or frequentist approaches?
- What is the modern view of the conflict (if any) between Bayesian and frequentist statistics?
- What are the best references (e.g., texts or pedagogical reviews) for introductory Bayesian statistics and for advanced topics?
 - As we compile lists: What are we missing? Are there more modern versions?
- What are the common or subtle pitfalls that novices to Bayesian methods fall into?
- What are we likely unaware of on the frontier of (Bayesian) statistical methods?
 - D. Furnstahl: In interacting with applied mathematicians I’ve found that physicists are often using the Numerical Recipes version of numerical methods, while the state-of-the-art is one or two generations more advanced. What are the analogs for statistics?
 - A Steiner: I’m currently using Goodman and Weare (2010)’s affine-invariant MCMC. Is there any way to do better? I’d like to get more accurate results with fewer samples. Will Metropolis-Hastings methods be superior if I have a sufficiently accurate proposal distribution?

1.3.2 Parameter estimation, model calibration, and model selection

- What is the difference between model calibration and parameter estimation?
- How should one do basic regression analysis?

- The old-school theoretical physics way is to do a least-squares fit with adding penalty terms for theoretical errors (which could be from the model or from the numerical method used to calculate the model) in quadrature to the data errors.
- When the theoretical systematic uncertainty is not known, one often determines the overall scale by requiring $\chi^2/\text{dof} = 1$ (Birge factor). How is this done in Bayesian statistics?
- When should a nuclear model with systematic theory errors have a statistical distribution of residuals?
- What are appropriate Bayesian priors?
- A. Steiner: How does one deal with the ambiguity created by heteroscedasticity? E.g. if we have two types of data points in a χ^2 fit, how do we decide the relative theoretical uncertainty between the two types?
- What approximations or techniques are useful for reducing computational cost?
- What is Approximate Bayesian Computation?
- What method should I use for calculating the evidence or odds ratios?
 - e.g., simulated annealing, nested sampling, analytic approximations, ...
 - What are the pros and cons?
- How do we propagate theoretical uncertainties (e.g., from truncations of an expansion or limitations of a physics model) to calculations of physics observables?

1.3.3 Priors

- What is Bayesian model checking and how can it be used to minimize or validate the influence of priors?
- What are other ways to validate priors?
- How does empirical Bayes work and when is it useful (or dangerous)?
- How do we choose priors for systematic errors in physics?
 - E.g., what general guidance is there?
 - What range of priors should I consider?
 - How does one choose a “non-informative” prior?

1.3.4 Software

- What should we know about MCMC sampling algorithms and software?
 - MCMC programs are often a black box to physicists.

- What are recommended implementations for different types of physics applications?
- Are there parallelized versions?
- What are the pitfalls or “tricks” in using MCMC?
- Should one use more than one algorithm?
- Autocorrelations in MCMC
 - * A. Steiner: I’m using the method outlined <http://www.math.nyu.edu/faculty/goodman/teaching/MonteCarlo2005/notes/MCMC.pdf> similar to the [acor](#) program used in [emcee](#).
- What are good programs for visualization (e.g., of projected posteriors)?
- What are the best software options for Python, C++, R, ...

1.3.5 Other topics

- Inconsistent data (or model)
- Outliers
- Model and uncertainty extrapolation
- Empirical Bayes
- Emulation
- A Steiner: In nuclear astrophysics, in order to perform a proper uncertainty quantification, we need two things: (i) the correlations between masses in the Atomic Mass Evaluation, and (ii) the correlations between parameters in popular mass models (e.g. FRDM). How do we get those?
- A. Steiner: What can be understood from the analogy between a particle propagator and a conditional probability distribution? Or does the fact that the former is defined over complex numbers spoil the analogy?

2 Background

[It is not yet clear how to organize the subsections, which currently overlap significantly.]

2.1 Basics of Bayesian statistics

- Not an exhaustive recounting, but the minimum needed to make sense of the review.
- Pointers to the literature (which is listed later).

2.2 Conceptual or philosophical issues

- In comparison to frequentist approaches, the emphasis is on sampling of posteriors rather than optimization (e.g., finding the maximum of the likelihood).
- Younger generation statisticians do not have the baggage of the frequentist-Bayesian wars; they freely use both approaches as tools to do statistical analysis.

3 Themes

3.1 Implication of Bayesian philosophy/practices/concepts

[*This may go in section 2.2 instead.*]

- emphasis of posterior vs. minimizing of χ^2
- sampling vs. optimization

3.2 Discrepant functions / random effects

3.3 Gaussian Process (GP) as a tool

- emulators
- discrepant functions / random effects
- linear regression \rightarrow GP

3.4 χ^2 and dofs

- χ^2/dof is not meaningful for Bayesian statistics
- dofs when there are priors
- how should one count dofs?
- augmented χ^2
- When does χ^2/dof make sense?

3.5 Pitfalls

- using same data to estimate priors and determine uncertainty

3.6 Model selection / comparison (or metrics)

- laundry list of approaches (see Vera's talk)
- new approach for nuclear physics: mixture models

3.7 Selection of priors

- uniform vs. where it makes a difference
- knowledge of underlying model
- running against the boundary

3.8 Use of MCMC

- must be able to evaluate quickly enough
- fast model vs. good emulator vs. slow function (dictates what you do)
- burn-in as usual; is checking autocorrelation and skipping critical?
- nested sampling?

3.9 Specific issues

- Instabilities in parameter space
 - undefined χ^2
 - diagnostics (e.g., response function)
 - alternatives?

3.10 Understanding of model function

- ???

A Guide to the literature (written and software)

- The vastness of the literature on statistical methods poses a problem for the nuclear physicist: there is not an adequate sampling algorithm. This is compounded by unfamiliar notation and terminology used by the statistics community, as well as conceptual differences in the formulation of problems.
- Here we propose a sampling of books and articles.

A.1 General Bayesian statistics

The following are often recommended as introductory guides for physicists, because the examples and language are drawn from physics:

- *Data Analysis: A Bayesian Tutorial* by Devinderjit Sivia and John Skilling [1].
- *Bayesian Logical Data Analysis for the Physical Sciences* by Phil Gregory [2].
- *Bayes in the sky: Bayesian inference and model selection in cosmology* by Roberto Trotta [3].

Recommended more advanced texts include:

- *Bayesian Data Analysis* by Gelman et al. [4].

A.2 Gaussian Process (GP) models

Software:

- The “[Sample Size Calculations for Computer Experiments](#)” app provides a sandbox for playing with Gaussian processes. Read the “About” tab first; it includes definitions of parameters used in the GP “correlation family” (which define the covariance). The “Sample Path Plots” tab has an app that lets you draw samples of functions that depend on user selected parameters (actually “hyperparameters”) that specify the details of the correlation family.
- Dan Foreman-Mackey’s [Gaussian Process Python Library and Tutorial](#) steps you through using GPs with [emcee](#).

References:

-

A.3 Hypothesis Testing and Model Comparison

B Summary of notation and terminology

B.1 Lexicon

hyperparameter This is ...

C Using git and github

[This section is just for building the white paper. It will not be included in the production version.]

References

- [1] D. Sivia and J. Skilling, *Data Analysis: A Bayesian Tutorial*. Oxford University Press, 2006. [11](#)
- [2] P. Gregory, *Bayesian Logical Data Analysis for the Physical Sciences*. Cambridge University Press, 2005. [11](#)
- [3] R. Trotta, *Bayes in the sky: Bayesian inference and model selection in cosmology*, *Contemp. Phys.* **49** (2008) 71–104, [arXiv:0803.4089 \[astro-ph\]](#). [11](#)
- [4] A. Gelman, J. B. Carlin, H. S. Stern, and D. B. Rubin, *Bayesian Data Analysis*. Chapman and Hall/CRC, 2003. [11](#)